

# Improper Request Propagation before Data Reception in Write Transactions in a Bus Architecture

Francesco Restuccia

UC San Diego

# High-level summary of the proposed weakness

The product uses a multi-manager shared bus architecture with an on-chip interconnect that speculatively propagates a request for a write transaction, but it does not ensure that the remaining steps required by the manager to complete the request are fulfilled in a timely manner.

**Key idea:** speculative forwarding of write address before data arrival lead to DoS/delays in the access of shared resources

- Shared port to resource becomes reserved; stalled data phase blocks others
- Turns local (IP) misbehavior into system-wide availability issue (DoS)
- Recovery might require system interconnect reset if not handled properly

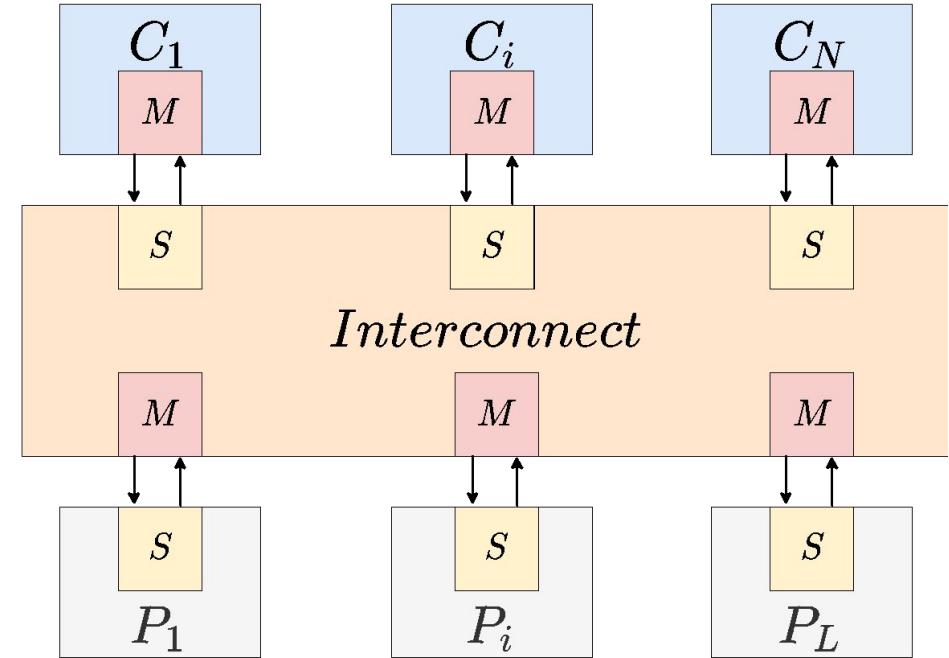
# Reference architecture

**Managers/Controllers (active):**

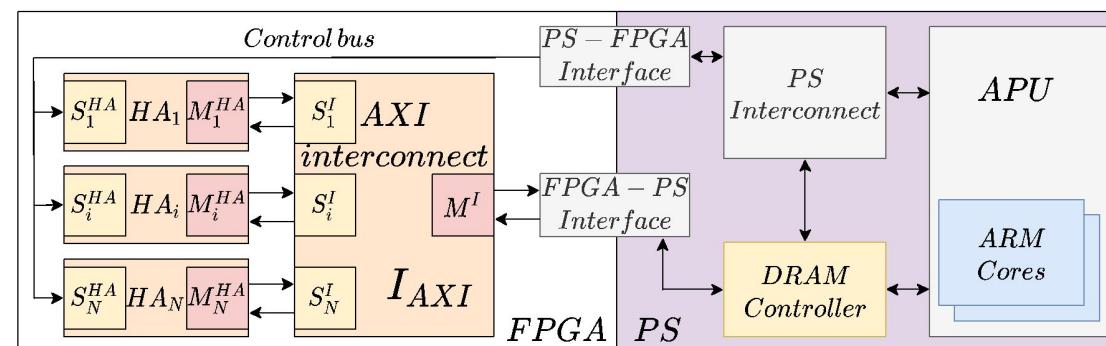
processors, DMAs, hardware accelerators,  
etc.

**Interconnect:** arbitrates the access of  
controllers to peripherals

**Peripheral (passive):** shared, e.g.,  
memories, I/Os, etc.



Anatomy of a typical heterogeneous platform  
(simplified)



A sample multi-manager architecture deployed on an FPGA SoC platform

# Minimal Threat Model and Assumptions

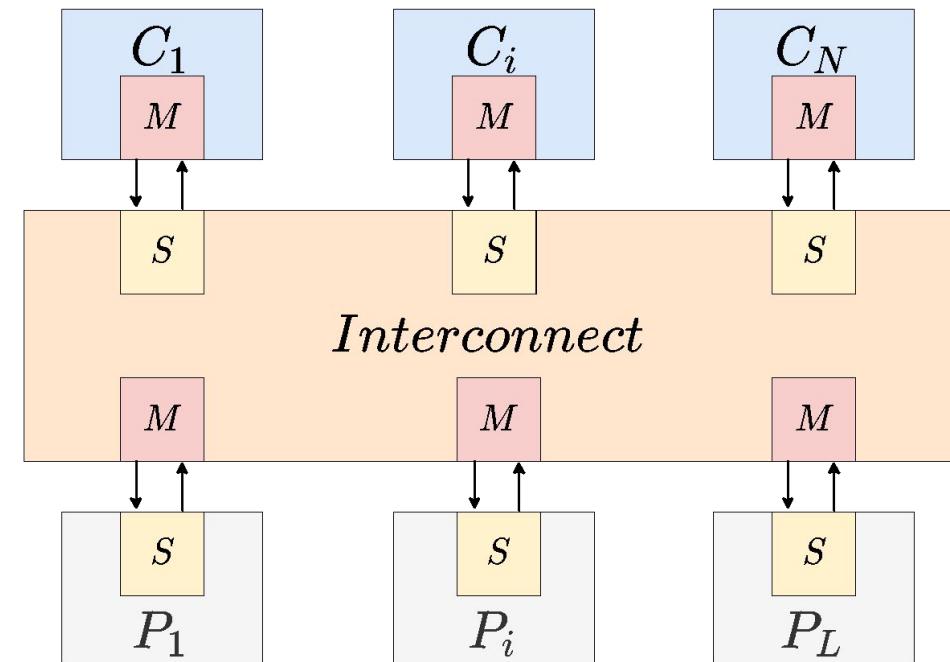
Multi-manager SoC interconnect (e.g., AXI) with shared resources.

- Interconnect assumes cooperative managers and timely data provision.

- Interconnect is trusted

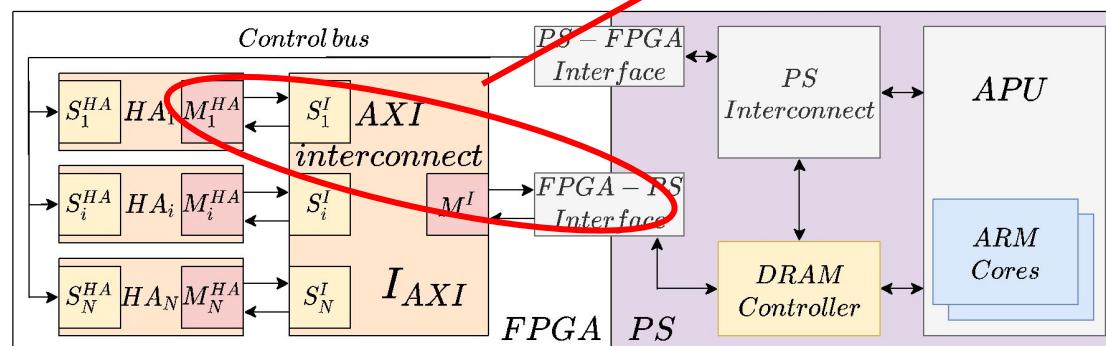
- Peripherals are trusted

- Managers might not be trusted
  - Malicious/faulty/misbehaving
  - Third-party IPs



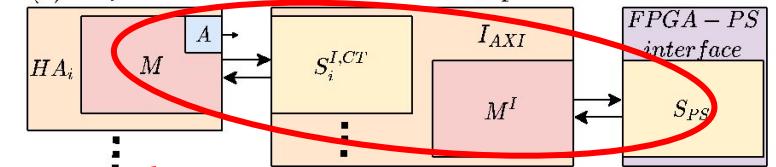
# Example

The interconnect receives a write request from a manager initiator (e.g., sending the header to the interconnect for writing N words of data, with N ranging from 1 to the maximum burst length allowed).



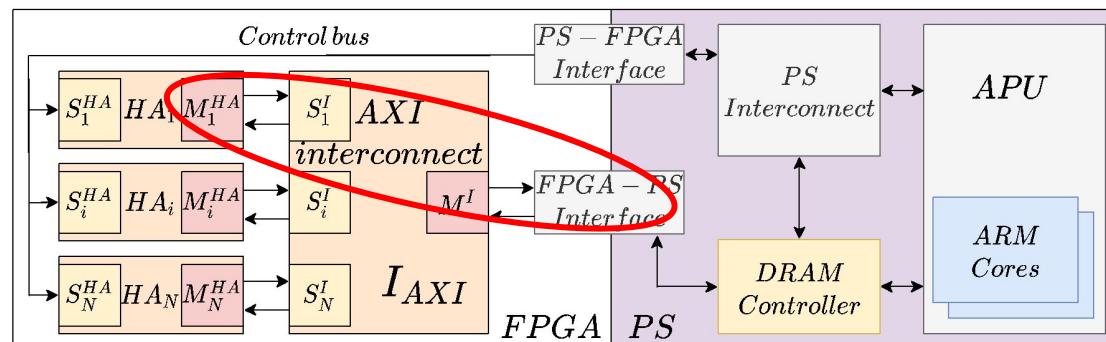
## Sample 4-word write transaction - CT switching

(a)  $HA_i$  issues the 4-word burst address request  $A$



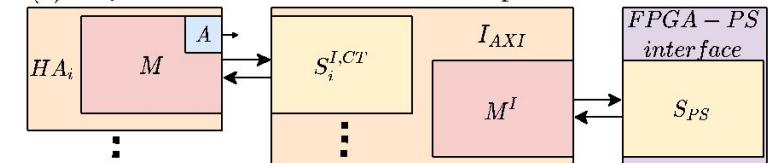
# Example

The interconnect then speculatively propagates the request to the target peripheral through the shared bus before the corresponding data are received at the interconnect.

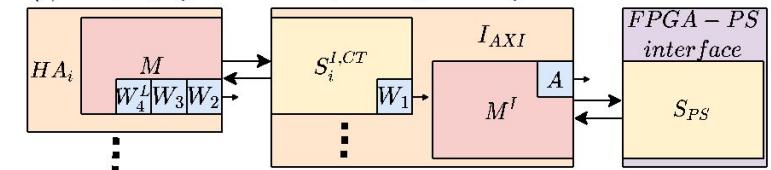


## Sample 4-word write transaction - CT switching

(a)  $HA_i$  issues the 4-word burst address request  $A$

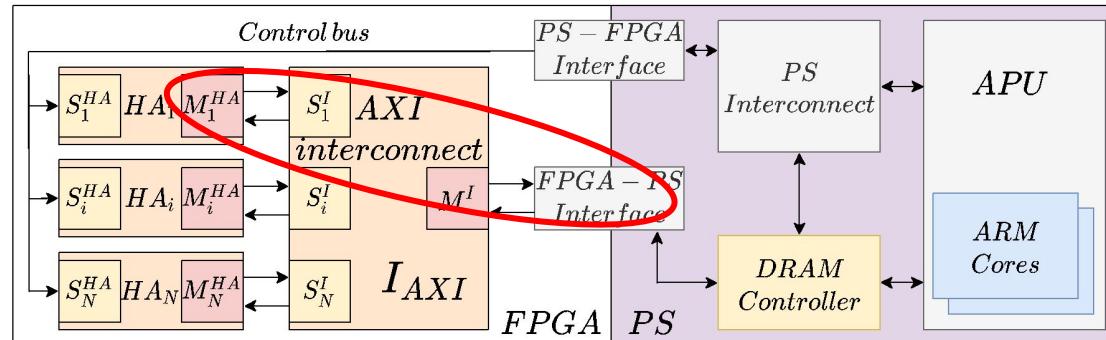


(b)  $A$  is propagated and  $HA_i$  starts providing the data words



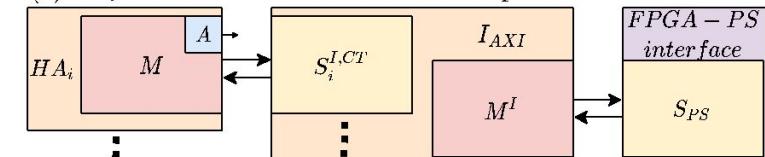
# Example

This speculatively reserves the bus to the manager transaction issuer, which is left free to delay its data provisioning phase. This reserves the whole bus path between the initiator and the target peripheral until all of the words associated with the request are provided by the manager.

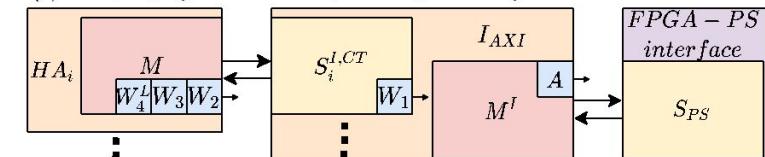


## Sample 4-word write transaction - CT switching

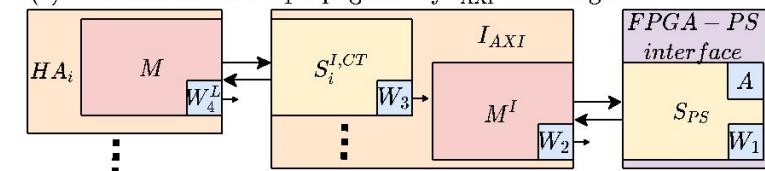
(a)  $HA_i$  issues the 4-word burst address request  $A$



(b)  $A$  is propagated and  $HA_i$  starts providing the data words



(c) The data words are propagated by  $I_{AXI}$  following  $A$



# Example

**How stalls are generated?**

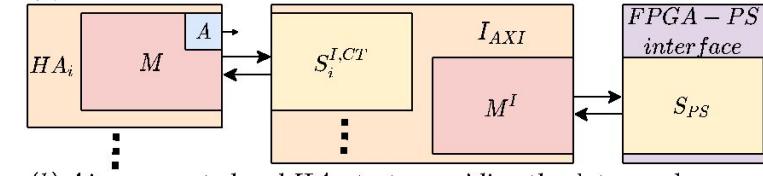
Typical switching method: **Cut-through**

+

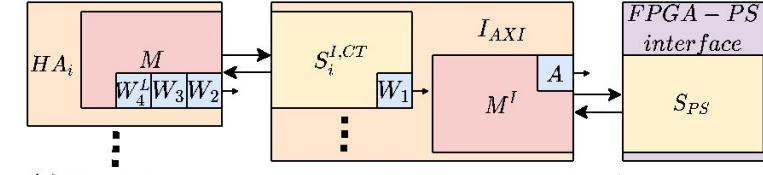
**Protocol definition:** write data must be propagated in the same order of the request for transactions

**Sample 4-word write transaction - CT switching**

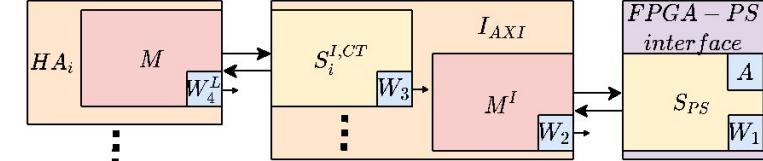
(a)  $HA_i$  issues the 4-word burst address request  $A$



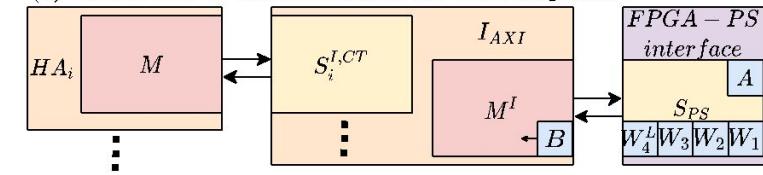
(b)  $A$  is propagated and  $HA_i$  starts providing the data words



(c) The data words are propagated by  $I_{AXI}$  following  $A$



(d) The FPGA-PS interface issues the write response  $B$



# Example

**How stalls are generated?**

Typical switching method: **Cut-through**

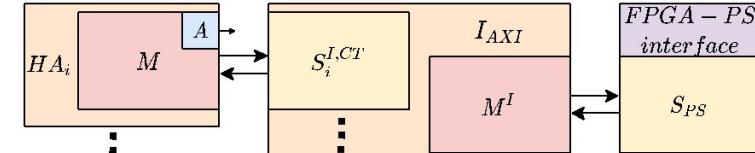
+

**Protocol definition:** write data must be propagated in the same order of the request for transactions

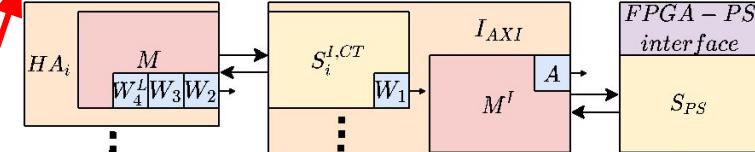
**(b) The hardware modules are trusted to provide (quickly) all of the data words after submitting a request for transaction**

**Sample 4-word write transaction - CT switching**

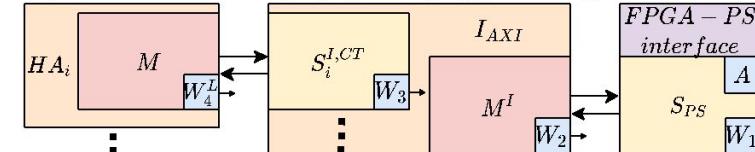
(a)  $HA_i$  issues the 4-word burst address request  $A$



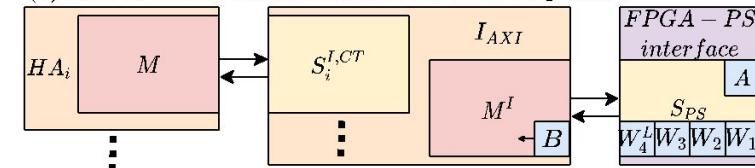
(b)  $A$  is propagated and  $HA_i$  starts providing the data words



(c) The data words are propagated by  $I_{AXI}$  following  $A$



(d) The  $FPGA-PS$  interface issues the write response  $B$



# Example

## How stalls are generated?

Typical switching method: **Cut-through**

+

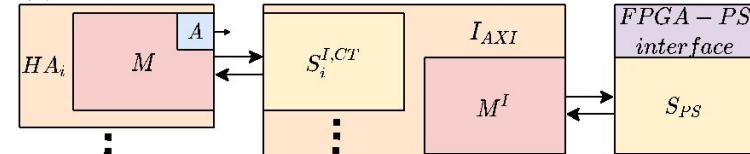
**Protocol definition:** write data must be propagated in the same order of the request for transactions

**(b) The hardware modules are trusted to provide (quickly) all of the data words after submitting a request for transaction**

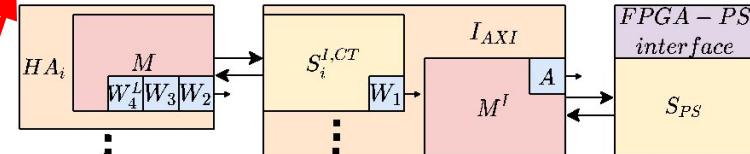
Example: AMBA AXI specification

### Sample 4-word write transaction - CT switching

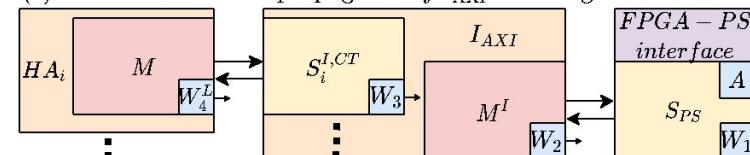
(a)  $HA_i$  issues the 4-word burst address request  $A$



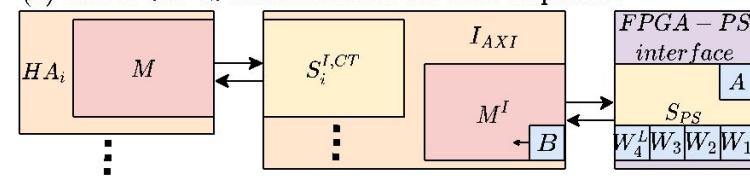
(b)  $A$  is propagated and  $HA_i$  starts providing the data words



(c) The data words are propagated by  $I_{AXI}$  following  $A$



(d) The FPGA-PS interface issues the write response  $B$



### A5.2.2 Write data ordering

A Manager must issue write data in the same order that it issues the transaction addresses.

An interconnect that combines write transactions from different Managers must ensure that it forwards the write data in address order.

The interleaving of write data with different IDs was permitted in AXI3, but is deprecated in AXI4 and later. See the AMBA AXI and ACE Protocol Specification issue F specification for more details on write data interleaving.

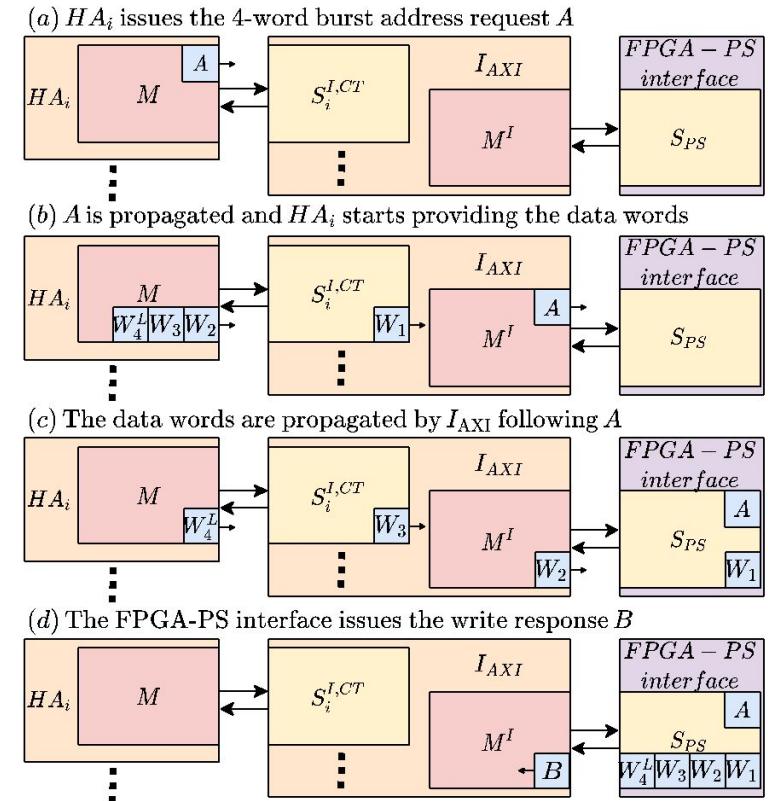
# Summary of Demonstrative Example

## Cut-Through Write

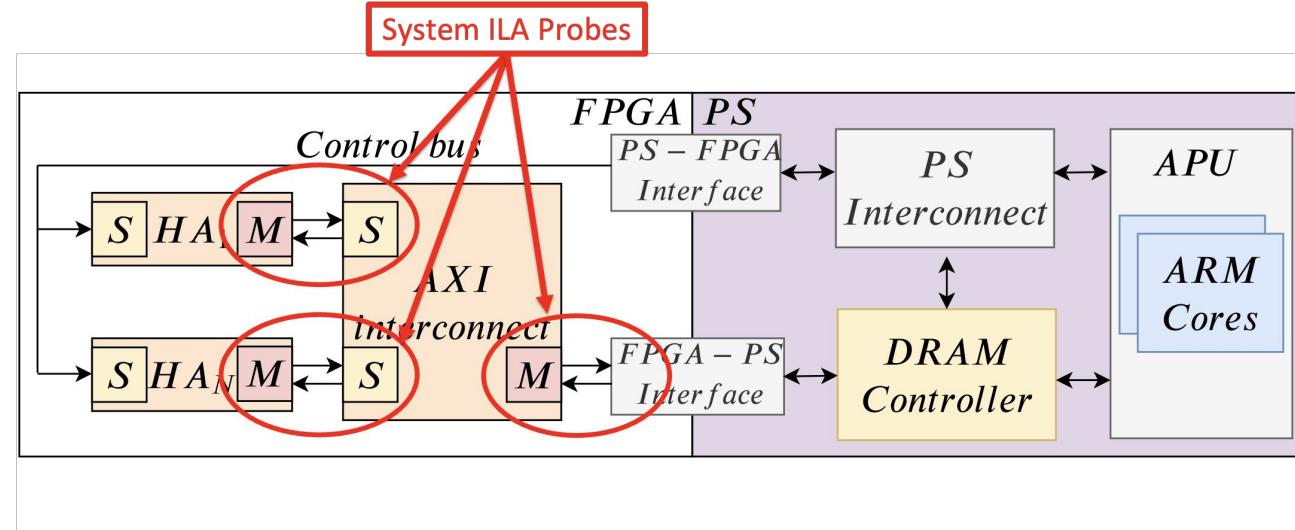
- Manager issues write address A (burst).
- Interconnect forwards A to shared port before data W arrive.
- Shared port becomes booked for A; other managers blocked.
- Manager delays W → prolonged stall; bus/path occupied.
- AXI lacks a native abort for such stalled transactions.

## Where a DoS is generated

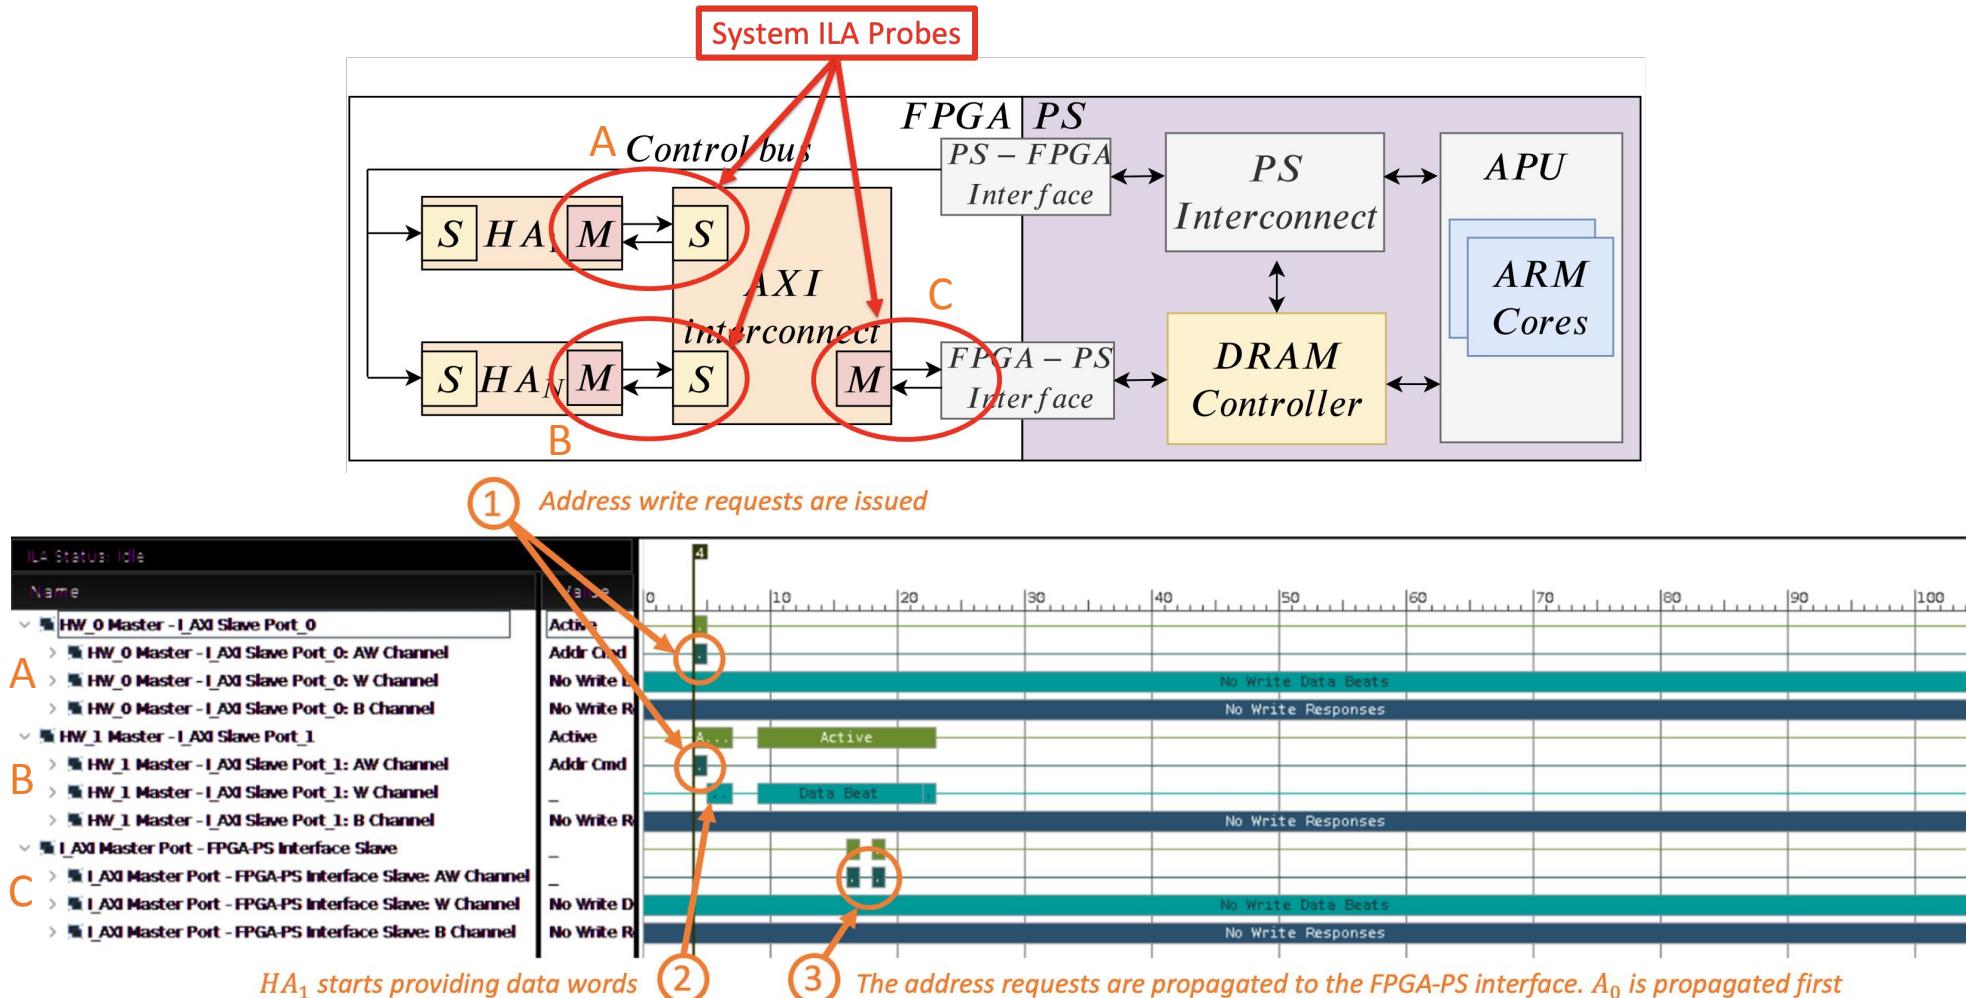
- Speculative reservation without bounded data timing.
- Single misbehaving manager blocks shared link for all.
- Stalls propagate through cut-through paths.
- Lack of protocol-level abort of bus transactions



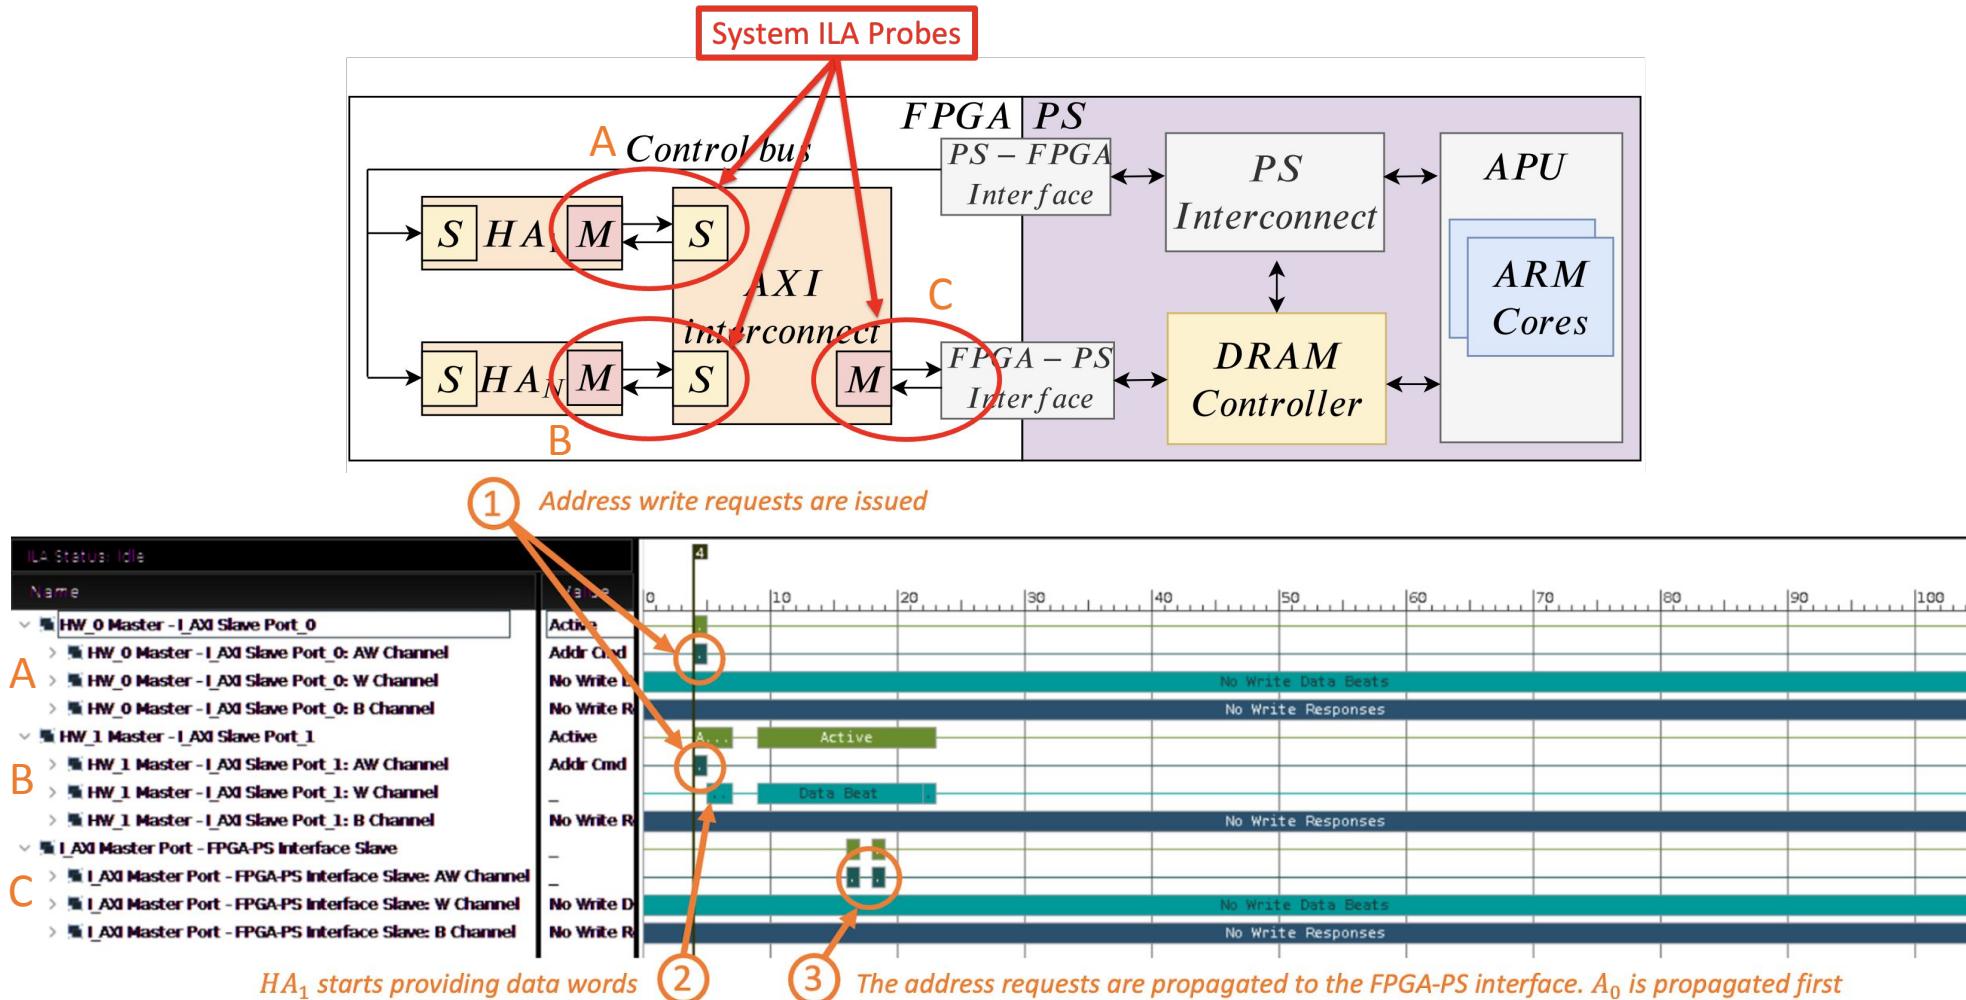
# Demonstration on COTS FPGA



# Demonstration on COTS FPGA



# Demonstration on COTS FPGA



One misbehaving module cause denial of service in accessing the shared DRAM memory

# Demonstration - commercial manager IP

Real waveform track collected from an execution of the Xilinx DPU core (v 3.3) on a ZYNQ Ultrascale+ platform (ZCU102)



# Demonstration - commercial manager IP



The transaction is  
issued

Partial data provided

The DPU implements this behaviour for speculative bus access. This behaviour is  
fully compliant with the protocol standard

# Demonstration - commercial manager IP



**The DPU implements this behaviour for speculative bus access. This behaviour is fully compliant with the protocol standard**

If the DPU shares the bus with other hardware modules, these last cannot access the shared peripheral while the bus is stalled by the DPU core

# Modes of Introduction

## Where it gets introduced

- Phase: Architecture & Design.
  - Unsafe switching (e.g., cut-through) with multi-channel standards.
  - Trust in manager modules; lack of bounded timing for data after address.
- Risk increases with third-party interconnect IP when switching policy is undocumented.

# Platforms & Technologies dependencies

- Languages/OS: Not language- or OS-specific.
- Architectures: SoC; examples include ARM-based interconnects.
- Technology: Bus/Interface, Communication, Network-on-Chip.
- Class: System-on-Chip; prevalence depends on interconnect choice.

AXI is a representative example, not exclusive.

# Impact

- Scope: Availability.
- Denial-of-Service on shared resources (memory, I/O).
- System-wide delays; missed deadlines in real-time domains.
- Possible need for interconnect reset; degraded reliability.

Implications particularly in critical systems (automotive, avionics, space, medical, robotics, etc.).

# Detection Methods

## Architecture/Design Review

- Switching policy (cut-through).
- Availability protection on shared resources.
- Availability protection on the untrusted managers.

## Verification

- Properties to prove bounded stalls to shared resources.
- Temporal properties for progress/completion.

## Simulation / Emulation

- Inject delayed/withheld data from manager side.
- Observe stall propagation in simulation.
- Some VIP may not trigger this (as it is compliant with the standard)

# Mitigations & Trade-offs

## Store-and-Forward

- Legacy solution: Buffer all data before booking shared resource.
- Prevents propagation of upstream stalls to the rest of the system.
- Trade-offs: added considerable latency and area.

## Cut-and-Forward

- Split bursts into bounded chunks (C).
- Book shared port per chunk only when data is already buffered.
- Balances latency/area vs availability; configurable C.
- Tradeoff: shorter transactions - might impact average performance

## Monitor & Recovery

- Detect long stalls at the manager or peripheral level.
- Abort/complete stalled transactions according to the standard when stalls are detected.
- Restore interconnect to usable state.
- Tradeoff: additional logic + analysis for configuration

# Observed Examples & Relations

## Observed Examples

- Observed in commercial AXI interconnect IPs
- Cut-through used as default in commercial interconnects (performance)

## Related Weaknesses

- ChildOf: CWE-841 (per submission).
- Cross-reference with availability/DoS-centric hardware CWEs.

## Recommendations & Checklists

## Guidance for Architects & Integrators

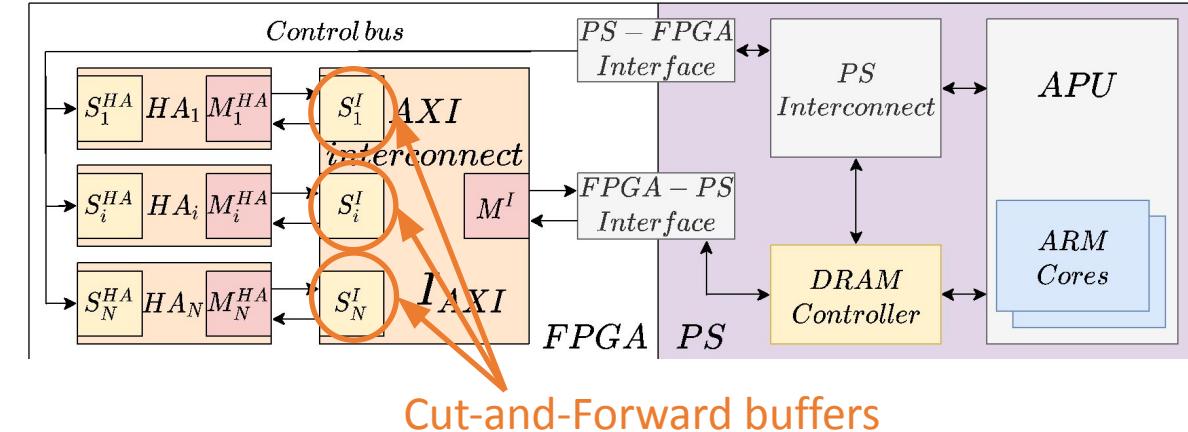
- Review switching policy from interconnect IP vendors.
- Enforce bounded timing from address to data on writes.
- Prefer S&F or C&F in critical scenarios or add monitors for recovery pathways.
- Verify IPs with properties and targeted stall tests.

# The Cut and Forward (C&F) switching

## What is Cut and Forward?

Novel **switching methodology** to be integrated in an AXI interconnect

AXI compliant – **fully transparent** to the **network** and resources (i.e., managers and peripherals)



## Why Cut and Forward?

Can combine **low-latency** transactions propagation and **low area footprint** (Cut-Through) with **safe and secure** bus communications **for any workload** (Store-and-Forward)

Configurable on a target application through the C&F parameter **C**

# How Cut and Forward works

Sample 4-word write transaction – C&F switching configured with C=2

C&F buffers are configured with the parameter C

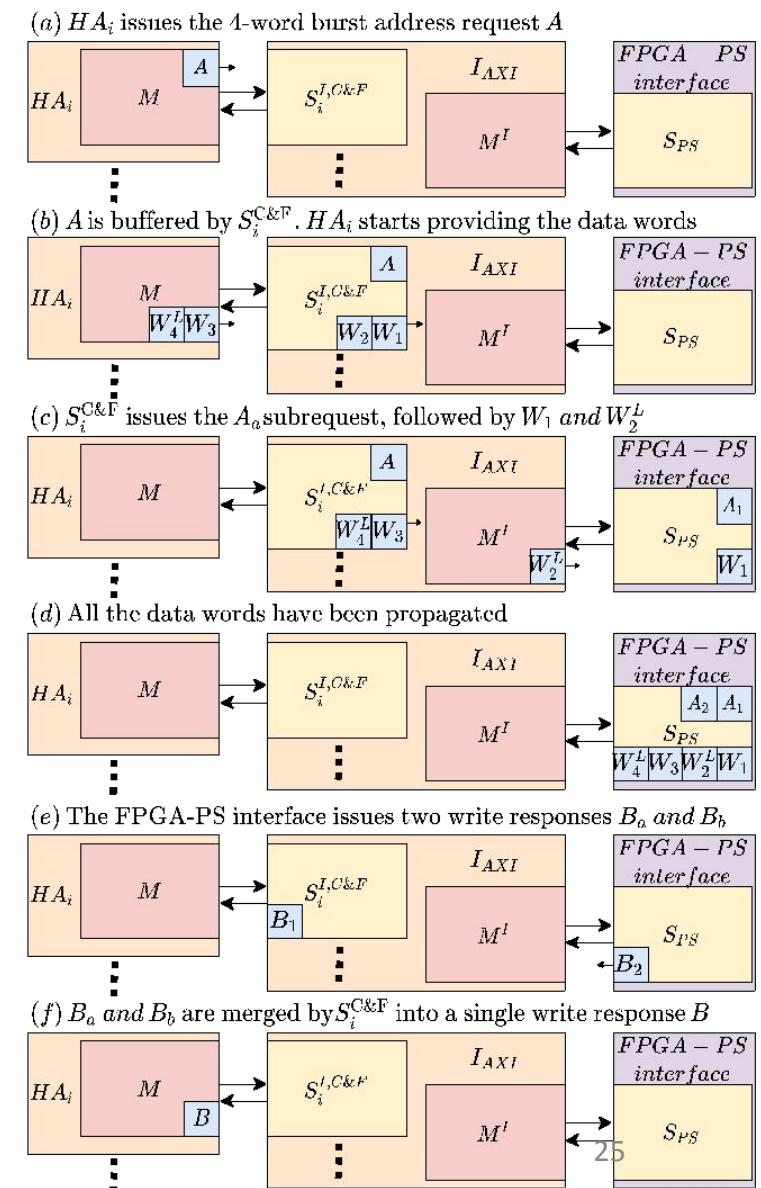
C&F waits for C data words before propagating any address request (i.e., booking the shared bus)

When C words have been buffered, C&F generates and sends a subrequest for C words (i.e., the bus is booked only for C words)

The buffered C words are following propagated by C&F with no stalls

In parallel with the propagation of the C words, the following C words can be sampled by C&F. This introduce a pipelining effect between data buffering (validation) and data propagation

The B responses are merged by C&F – compliant with the AXI standard



# The Carfield SoC



Journals & Magazines > IEEE Transactions on Circuits... > Early Access [?](#)

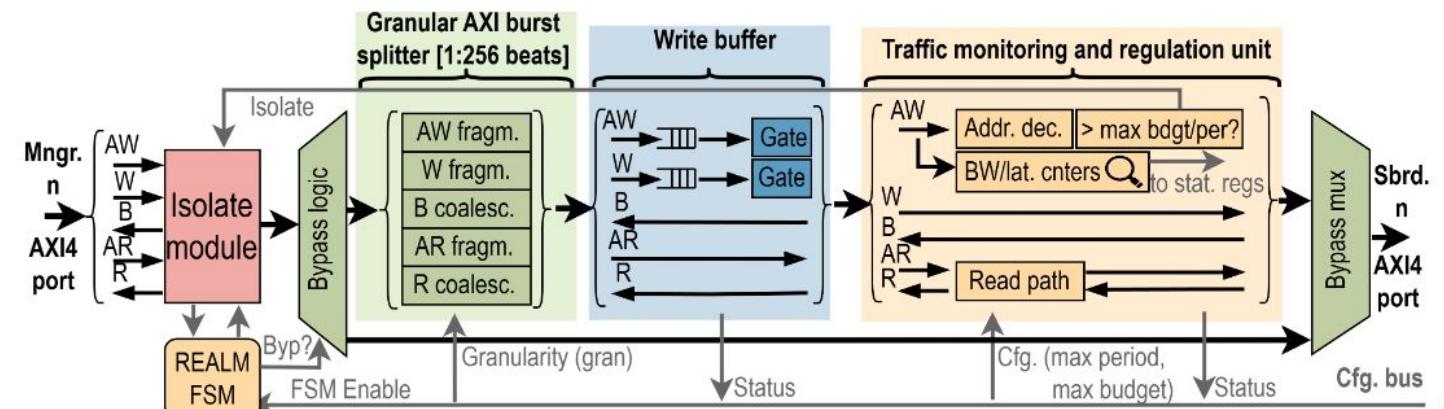
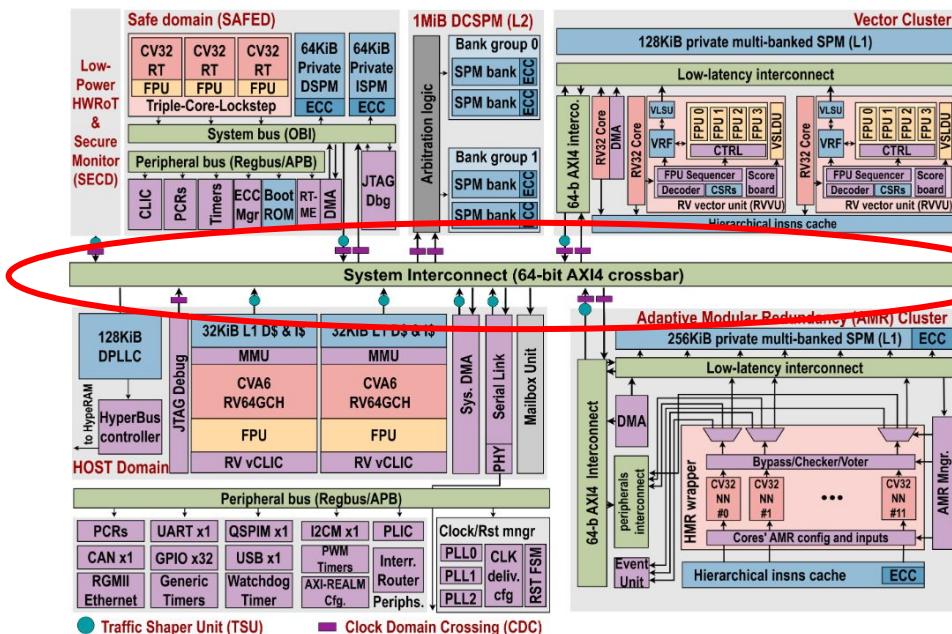
## A Reliable, Time-Predictable Heterogeneous SoC for AI-Enhanced Mixed-Criticality Edge Applications

Publisher: IEEE

Cite This



Angelo Garofalo · Alessandro Ottaviano · Matteo Perotti ; Thomas Benz ; Yvan Tortorella ; Robert Balas All Authors



A. Garofalo *et al.*, "A Reliable, Time-Predictable Heterogeneous SoC for AI-Enhanced Mixed-Criticality Edge Applications," in *IEEE Transactions on Circuits and Systems II: Express Briefs*, doi: 10.1109/TCSII.2025.3591225. keywords: {Hardware;Vectors;Artificial intelligence;Software;Bandwidth;Logic;Energy efficiency;Electric breakdown;Safety;Runtime;Heterogeneous mixed-critical SoC;time-predictable SoC;hardware acceleration;fault tolerant hardware},

# References

- F. Restuccia and R. Kastner, "Cut and Forward: Safe and Secure Communication for FPGA System on Chips," in IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems, vol. 41.
- F. Restuccia, A. Biondi, M. Marinoni and G. Buttazzo, "Safely Preventing Unbounded Delays During Bus Transactions in FPGA-based SoC," 2020 IEEE 28th Annual International Symposium on Field-Programmable Custom Computing Machines (FCCM), Fayetteville, AR, USA
- Meza, Andres, et al. "Safety verification of third-party hardware modules via information flow tracking." Proc. 1st Real-Time Intell. Edge Comput. Workshop (RAGE) Co-Located 59th Design Autom. Conf.(DAC).
- T. Benz et al., "AXI-REALM: Safe, Modular and Lightweight Traffic Monitoring and Regulation for Heterogeneous Mixed-Criticality Systems," in IEEE Transactions on Computers.
- A. Garofalo et al., "A Reliable, Time-Predictable Heterogeneous SoC for AI-Enhanced Mixed-Criticality Edge Applications," in IEEE Transactions on Circuits and Systems II: Express Briefs.

frestuccia@ucsd.edu